

Cerebras Systems

AI Supercomputing / Wafer-Scale Compute · Sunnyvale, CA · Nasdaq: CBRS (blockbuster IPO May 2026; ~\$67B cap; \$510M 2025 rev, profitable)

Cerebras - the wafer-scale AI-supercompute leader, fresh off a blockbuster ~\$67B Nasdaq IPO - has never been in motorsport, and for a greenfield F1 team building its simulation stack it is a real compute workstream and a two-way exchange, not a logo.

79

OPPORTUNITY / 100

TIMING WINDOW

HOT

SERIES

F1

RECOMMENDED TEAM

Cadillac

ACTION HORIZON

6-10 WKS

THE CASE

Cerebras Systems, led by co-founder & CEO Andrew Feldman, builds wafer-scale AI processors (the WSE-3 is the largest chip ever made, ~4 trillion transistors) and the systems around them, and IPO'd on Nasdaq (CBRS) in a blockbuster May 2026 - ~\$67B market cap, \$510M 2025 revenue at a 47% net margin (profitable), with multi-billion-dollar compute agreements with OpenAI and AWS. It has never touched motorsport. The F1 fit is real: aerodynamics is the biggest lap-time lever and it is compute-bound, so wafer-scale compute lets a team extract more CFD fidelity and faster iteration - and run power-unit, strategy and ML simulation - within F1's regulated CFD/wind-tunnel allocation (the ATR). HONEST framing: the ATR caps aero-CFD volume (so this maximises value within the cap rather than removing it), CAE software incumbents (Ansys/Cadence/Siemens) already occupy the toolchain (Cerebras is the compute underneath, a distinct sub-lane), and Cerebras carries real customer concentration (G42/MBZUAI was ~62% of 2025 revenue) - all noted, and the score reflects them. Decision-maker: Co-Founder & CEO Andrew Feldman.

WHY NOW

Cerebras is at a live capital + identity peak (blockbuster May-2026 IPO, ~\$67B cap) - exactly when a company invests in brand and lands marquee references - and F1's 2026 rules raise the simulation burden, so a founding compute relationship with a greenfield team is best built now. The honest why-now is catalyst-driven, not a dated race trigger.

PROOF POINTS — EVERY FIGURE VERIFIED TO A PRIMARY SOURCE (SEE P.2)

~\$67B

Blockbuster Nasdaq IPO (May 2026; priced \$185, opened +108%)

\$510M

2025 revenue (47% net margin; profitable)

WSE-3

Largest AI chip ever (WSE-3; ~4 trillion transistors)

Andrew Feldman

Co-Founder & CEO (since 2015; ex-SeaMicro, sold to AMD)

MBZUAI ~62%

Customer concentration (2025 revenue; G42-linked)

GRID FIT — CATEGORY-WHITESPACE CHECK ACROSS THE F1 GRID (1440 TEAM-FIT ENGINE)

Cadillac	PRIME LANE	greenfield entry — partner roster being built from zero
Red Bull Racing	CROWDED	Oracle
McLaren	OPEN	no rival in this category lane
Williams	OPEN	no rival in this category lane

BOTTOM LINE

The wafer-scale AI-supercompute leader (WSE-3, the largest AI chip ever) that IPO'd on Nasdaq in a blockbuster May 2026 (~\$67B cap; \$510M 2025 revenue at a 47% net margin - profitable; OpenAI + AWS multi-billion deals). REAL F1 workstream: aerodynamics is the #1 lap-time lever and it is compute-bound; wafer-scale compute lets a team extract more CFD fidelity and faster iteration within F1's regulated CFD/wind-tunnel cap (ATR), plus PU/strategy/ML simulation. Heavily deal-capable. HONEST caveats: F1's ATR caps aero-CFD volume, and Cerebras has customer concentration (G42/MBZUAI) - noted. Co-Founder & CEO Andrew Feldman.

WHY CADILLAC

Cadillac is the new American works team building its entire simulation and compute stack from zero, so an AI-supercompute partner can be the baseline it designs around (a genuine prime lane), and an American AI-compute champion alongside the American flag-bearer is a clean brand alignment - with GM behind it as a deep-pocketed, deal-capable counterpart. No AI-compute incumbent to conflict with (IFS runs ERP, TWG AI is data/analytics services, not wafer-scale HPC). It is a signal, not a placement - Cadillac's greenfield build is simply the best fit for a foundational compute partner.

VALUE TO CADILLAC — MODE A (TECH IN THE CAR / CHAMPIONSHIP)

A TWO-WAY exchange (tech + business both directions), not a logo fee - three concrete swaps. (1) COMPUTE <-> ELITE CFD BENCHMARK: Cerebras gives the team wafer-scale AI-compute to maximise CFD fidelity and iteration within the ATR cap (plus power-unit, strategy and ML simulation); in return F1 aerodynamics - one of the hardest real-world CFD/HPC problems - becomes Cerebras's credibility-defining benchmark and showcase to industrial and scientific buyers. (2) CO-ENGINEERED WORKLOADS <-> PRODUCT VALIDATION: jointly optimising the team's CFD/simulation workloads on wafer-scale hardware advances Cerebras's real-world HPC performance and reference architecture; the team gets faster, higher-fidelity aero development that converts directly to lap time. (3) TALENT + BRAND <-> ENTERPRISE REACH: an American compute champion on the American works team is a recruiting and brand halo for the HPC/ML engineers both sides compete for; the paddock's industrial/OEM ecosystem opens the enterprise HPC buyers Cerebras is scaling into. Cadillac gets a foundational simulation-compute edge from day one; Cerebras gets the hardest benchmark, a global stage and buyers.

DEAL ARCHITECTURE

Official AI Compute / Simulation Partner on Cadillac, minimum three years, structured as a compute-supply + co-engineering workstream (CFD/PU/strategy simulation on wafer-scale systems, joint HPC optimisation, a shared benchmark/reference programme) - value exchanged both ways, not a pure logo fee. Cerebras is highly deal-capable, so size to a genuine technical partnership.

PRIMARY DECISION-MAKER

Andrew Feldman

CO-FOUNDER & CEO, CEREBRAS SYSTEMS - VERIFIED

Co-founder and CEO of Cerebras since 2015 (prior: co-founded SeaMicro, sold to AMD), the owner of brand and strategic-partnership decisions and the entry point - for whom the hardest real-world simulation benchmark plus an enterprise-buyer stage is strategically valuable. Confirm the marketing/partnerships lead before outreach.

OPENING ANGLE

"Andrew - F1 aero is one of the hardest CFD problems on earth and it is compute-bound, and Cadillac is building its simulation stack from zero. Wafer-scale compute is their edge and their car is your benchmark and showcase. Fresh off your IPO, that's a two-way deal. 25 minutes?"

SCORE COMPOSITION

TIMING

18 / 20

Blockbuster Nasdaq IPO in May 2026 (priced \$185, opened +108%, ~\$67B cap) - a live, this-quarter capital + identity inflection at the peak of the AI-compute cycle.

CAPACITY

18 / 20

~\$67B market cap, \$510M 2025 revenue at a 47% net margin (profitable), OpenAI + AWS multi-billion deals - deal-capable at any tier (the standout, and the exact capacity smaller compute names lacked).

BRAND FIT

15 / 20

AI supercompute maps onto F1's compute-bound aerodynamics (the #1 performance lever) and a cutting-edge engineering story - strong, though CAE software incumbents (Ansys/Cadence/Siemens) already sit in the CFD toolchain, so it is a distinct compute-hardware sub-lane.

URGENCY

13 / 20

Fresh IPO plus rising F1-2026 simulation demand, but F1's Aerodynamic Testing Restriction caps CFD volume and there is no dated Cerebras-specific trigger.

OPS FIT

15 / 20

MODE A: wafer-scale compute genuinely accelerates a team's CFD, power-unit and strategy/ML simulation - higher fidelity and faster iteration per ATR-allocated unit, and broader simulation beyond capped aero-CFD; a real workstream, honestly bounded by the ATR cap.

RISKS & COUNTERS

F1'S CFD CAP (ATR) F1 regulates CFD/wind-tunnel volume per team (Aerodynamic Testing Restriction), so compute can't simply be scaled up without limit. **Counter:** Wafer-scale value is higher fidelity and faster iteration PER allocated unit, plus uncapped power-unit/strategy/ML simulation - maximising the regulated budget, not evading it.

CAE SOFTWARE INCUMBENTS Ansys, Cadence and Siemens already sit in the CFD toolchain. **Counter:** Cerebras is the compute layer underneath that software - a distinct hardware sub-lane; position as accelerating the existing toolchain, not replacing it.

CUSTOMER CONCENTRATION / INDEPENDENCE G42/MBZUAI was ~62% of 2025 revenue and G42 is Cerebras's largest backer (a concentration and geopolitical-scrutiny risk). **Counter:** Revenue is diversifying fast (OpenAI, AWS multi-billion deals) and the company is profitable at ~\$67B - deal-capability is not in question; note the concentration honestly.

SOURCES

- <https://finance.yahoo.com/markets/stocks/articles/cerebras-raises-5-55-billion-003314885.html>
- <https://tech-insider.org/cerebras-ipo-filing-510m-revenue-openai-deal-23b-valuation-2026/>
- <https://www.cerebras.ai/press-release/cerebras-announces-third-generation-wafer-scale-engine>
- <https://en.wikipedia.org/wiki/Cerebras>
- <https://www.agbi.com/tech/2026/05/g42-linked-chipmaker-cerebras-jumps-in-first-day-on-nasdaq/>